



Boosting RDKit molecular simulations through OpenMM

Paolo Tosco

Outline

> Background



> Implementation



> Results



> Conclusions and outlook

Background

Background



> OpenMM

- > High performance toolkit for molecular simulation
- > Python, C, C++, Fortran (!) bindings
- > Open source
- > Actively maintained on GitHub
- > Licensed under MIT and LGPL

> <http://openmm.org>

OpenMM

About OpenMM

Backed by researchers and developers from Stanford University, MSKCC, and others around the world.



Custom Forces

Want a custom force between two atoms? No problem. Write your force expressions in string format, and OpenMM will generate blazing fast code to do just that. No more hand-writing GPU kernels.

Highly Optimized

OpenMM is optimized for the latest generation of compute hardware, including AMD (via OpenCL) and NVIDIA (via CUDA) GPUs. We also heavily optimize for CPUs using intrinsics.

Portable

We strive to make our binaries as portable as possible. We've tested OpenMM on many flavors of Linux, OS X, and even Windows.

What about...



OpenMM-accelerated
Open-Source Cheminformatics
and Machine Learning

> ...spicing up the RDKit MM simulations with OpenMM?

- > Licenses are compatible
- > So are APIs
- > Initial proof-of-concept on MMFF94
- > If worthwhile, extend to UFF

How much work is it?



OpenMM-accelerated
Open-Source Cheminformatics
and Machine Learning

- > OpenMM does not implement MMFF94 natively
 - > We need to implement it ourselves
 - > Can we implement it using Custom Forces?
 - > Can we borrow some Forces from already implemented force-fields?
 - > Fortunately, atom types and force constants are already assigned by the RDKit

Implementation

OpenMM side, C++

OpenMM implementation of MMFF94 Forces

$E_{\text{MMFF}} = \sum E_{\text{B}}_{ij}$	bond stretching	$E_{\text{B}}_{ij} = 143.9325 \frac{k b_{ij}}{2} \Delta r_{ij}^2 \left(1 + c s \Delta r_{ij} + \frac{7}{12} c s^2 \Delta r_{ij}^2 \right)$	OpenMM::CustomBondForce
$+ \sum E_{\text{A}}_{ijk}$	angle bending	$E_{\text{A}}_{ijk} = 0.043844 \frac{k a_{i,j,k}}{2} \Delta \vartheta_{ijk}^2 (1 + c b \Delta \vartheta_{ijk})$	OpenMM::CustomAngleForce
$+ \sum E_{\text{BA}}_{ijk}$	stretch-bend	$E_{\text{BA}}_{ijk} = 2.51210 (k b a_{i,j,k} \Delta r_{ij} + k b a_{i,k,j} \Delta r_{kj}) \Delta \vartheta_{ijk}$	OpenMM::AmoebaStretchBendForce
$+ \sum E_{\text{OOP}}_{ijk,l}$	out-of-plane bending	$E_{\text{OOP}}_{ijk,l} = 0.043844 \frac{k o o p_{i,j,k,l}}{2} \chi_{ijk,l}^2$	OpenMM::AmoebaOutOfPlaneBendForce
$+ \sum E_{\text{T}}_{ijkl}$	torsion	$E_{\text{T}}_{ijkl} = 0.5 [V_1 (1 + \cos \phi) + V_2 (1 - \cos 2\phi) + V_3 (1 + \cos 3\phi)]$	OpenMM::CustomTorsionForce
$+ \sum E_{\text{vdW}}_{ij}$	van der Waals	$E_{\text{vdW}}_{ij} = \varepsilon_{ij} \left(\frac{1.07 R_{ij}^*}{R_{ij} + 0.07 R_{ij}^*} \right)^7 \left(\frac{1.12 R_{ij}^{*7}}{R_{ij}^7 + 0.12 R_{ij}^{*7}} - 2 \right)$	OpenMM::AmoebaVdwForce
$+ \sum E_{\text{Q}}_{ij}$	electrostatic	$E_{\text{Q}}_{ij} = 332.0716 \frac{q_i q_j}{D (R_{ij} + \delta)^n}$	OpenMM::CustomNonbondedForce

The devil is in the details

OpenMM side, C++/C

The AMOEBA vdW term

OpenMM: :AmoebaVdwForce

$$\text{EvdW}_{ij} = \varepsilon_{IJ} \left(\frac{1.07 R_{IJ}^*}{R_{ij} + 0.07 R_{IJ}^*} \right)^7 \left(\frac{1.12 R_{IJ}^{*7}}{R_{ij}^7 + 0.12 R_{IJ}^{*7}} - 2 \right)$$

Unfortunately, none of the supported combination rules for R_{IJ}^* and ε_{IJ} applies to MMFF94

- > R_{IJ} is the distance between particles I and J
- > R_{IJ}^* is a combination of particle-specific R_I^* , R_J^* constants
- > ε_{IJ} is a combination of particle-specific ε_I , ε_J constants
 - > OpenMM supports three different combination rules for R_{IJ}^* and four for ε_{IJ}

The MMFF94 vdW term

MMFF94 combination rules are complex

$$\text{EvdW}_{ij} = \varepsilon_{IJ} \left(\frac{1.07 R_{IJ}^*}{R_{ij} + 0.07 R_{IJ}^*} \right)^7 \left(\frac{1.12 R_{IJ}^{*7}}{R_{ij}^7 + 0.12 R_{IJ}^{*7}} - 2 \right)$$

$$R_{IJ}^* = 0.5(R_{II}^* + R_{JJ}^*) \left(1 - B(1 - e^{-12\gamma_{IJ}^2}) \right) \quad (1) \quad \gamma_{IJ} = \frac{R_{II}^* - R_{JJ}^*}{R_{II}^* + R_{JJ}^*}$$

$$\varepsilon_{IJ} = \frac{181.16 G_I G_J \alpha_I \alpha_J}{\sqrt{\frac{\alpha_I}{N_I}} + \sqrt{\frac{\alpha_J}{N_J}}} \cdot \frac{1}{R_{IJ}^{*6}} \quad (2)$$

- > HBA/HBD particle pairs need their R_{IJ}^* and ε_{IJ} to be scaled down
- > In equation (1), B is 0 if one of the particles is a HBD, otherwise it is 0.2

- > The dependence of ε_{IJ} from R_{IJ}^* , and the need for scaling would make **OpenMM::CustomNonbondedForces** computationally inefficient

(AMOEBA + MMFF94)^{hurry} = BIG HACK!

- > I hacked the `OpenMM::AmoebaVdwForce` implementation adding a “MMFF” combination rule
- > I abused the per-particle σ_i , ε_i and `reduction` (not used by MMFF94) parameters passed to the `OpenMM::AmoebaVdwForce::addParticle()` method...
- > ...to pass R_i , $G_i a_i$ and a_i / N_i instead, using sign combinations to encode HBA/HBD features
- > I modified the platform-specific kernels accordingly

Yuk! But it works!

The MMFF94 electrostatic term

$$EQ_{IJ} = s \cdot k \frac{q_I q_J}{D(R_{IJ} + \delta)^n}$$

> The problem here is that the scaling factor **s** has to be **0.75** for 1-4 interactions, and **1.0** for 1-*n* interactions (*n* > 4)

1-4 interactions

1-*n* interactions
(*n* > 4)

> Unfortunately, only per-particle parameters are allowed, so the electrostatic term is split in two

> Per-particle and group exclusions do the rest

Implementation

RDKit side, C++

The `OpenMMForceField` class, or where the magic lies

The core of the new OpenMM-powered force field implementation is the `ForceFields::OpenMMForceField` class

```
class OpenMMForceField : public ForceField {
public:
    [...]
    OpenMM::System *getSystem() const;
    OpenMM::Context *getContext() const;
    void setIntegrator(OpenMM::Integrator *integrator);
    OpenMM::Integrator *getIntegrator() const
    void initializeContext();
    void initializeContext(const std::string& platformName, const std::map<std::string, std::string> &prop);
    void initializeContext(OpenMM::Platform& platform, const std::map<std::string, std::string> &prop);
    double calcEnergy(std::vector<double> *contribs = NULL) const;
    double calcEnergy(double *pos);
    void calcGrad(double *forces) const;
    void calcGrad(double *pos, double *forces);
    int minimize(unsigned int maxIts = 200, double forceTol = 1e-4, double energyTol = 1e-6);
protected:
    [...]
private:
    [...]
}
```

Setters/getters for `OpenMM::System`, `Platform` and `Integrator`

OpenMM-enabled re-implementations of the base class methods

Before people get nervous: support for OpenMM
can be *optionally* built into the RDKit 😊

The `MMFF::OpenMMForceField` class, where more magic lies

The MMFF94-specific machinery is hosted by the `MMFF::OpenMMForceField` class

```
class OpenMMForceField : public ForceField {
public:
    [...]
    void addBondStretchContrib(...);
    void addAngleBendContrib(...);
    void addStretchBendContrib(...);
    void addTorsionAngleContrib(...);
    void addOopBendContrib(...);
    void addVdWContrib(...);
    void addEleContrib(...);
    void addEleContrib1_4(...);
    const std::vector<std::string>& loadedPlugins();
    const std::vector<std::string>& failedPlugins();
protected:
    [...]
private:
    [...]
}
```

Before people get *even more* nervous: I haven't touched the base `ForceField` class, so ABI compatibility is guaranteed 😊

Get me an OpenMM-enabled force field, now!

To construct an OpenMM-enabled force field, all you need is call the familiar

```
OpenMMForceField *constructOpenMMForceField(ROMol &mol,  
    MMFFMolProperties *mmffMolProperties, double nonBondedThresh = 100.0,  
    int confId = -1, bool ignoreInterfragInteractions = true);
```

The only difference from the well-known call are those two tiny **OpenMM** prefixes

> Once you have created it, you may do the usual things:

> Calculate the potential energy: `ff->calcEnergy()`

> Run a minimization: `ff->minimize()`

> Or more exotic ones:

> Run *n* steps of MD: `ff->dynamics(n)`

Implementation

RDKit side, Python

I already know how to do this, don't I?

Also in Python, you shouldn't be too surprised by the new API:

```
MMFFGetMoleculeOpenMMForceField( (Mol)mol, (object)pyMMFFMolProperties [,
    (float)nonBondedThresh = 100.0 [, (int)confid = -1 [,
    (bool)ignoreInterfragInteractions = True]]])
```

Again, you'll only need to add a tiny **OpenMM** to your existing scripts

> Once you have created it, you may do the usual things:

> Calculate the potential energy: **ff.CalcEnergy()**

> Run a minimization: **ff.Minimize()**

> Or more exotic ones:

> Run *n* steps of MD: **ff.Dynamics(n)**

Results

On to the Jupyter notebook

Conclusions and outlook

A.k.a., the to-do list

Conclusions

- > The OpenMM implementation of MMFF94 within the RDKit
 - > delivers impressive performance even on consumer GPU hardware
 - > enables fast molecular mechanics simulations on both CPUs and GPUs
 - > Can be accessed with minimal modifications to old scripts

Outlook

- > Before I can get all this to you I need:
 - > To properly implement MMFF94 in OpenMM (no hacks!)
 - > To extend the implementation to UFF
 - > To add more APIs specific to molecular dynamics
 - > To put some thought in the API design to give full access to present and future OpenMM functionality



smarter chemistry | smarter decisions

Thank you for your attention

paolo@cresset-group.com

Acknowledgments

- > The OpenMM team for such a great open-source simulation toolkit
- > The RDKit team for... well, you know
- > Cresset for supporting this project



cressetgroup